

Kubernetes - By P Nageswara Rao

Gamut Gurus Technologies

&

Wiculy Learning Solutions

Email: info@Gamutgurus.com, info@wiculy.com

Contact: +91- 97393 68768

What is Kubernetes?

Kubernetes is:

- Kubernetes is an Orchestrator for containerized Applications
- Kubernetes is a Data centre OS

Docker & Kubernetes Integration

What is Docker & Kubernetes?

Docker: You write your application code in your favourite language, Build it, Test it, then use Docker to package it and ship it. *Yes... That's what Docker is.*


Then, who will handle the final step of running it on Test or Prod?

Hello Kubernetes 

Yes, It's Kubernetes! Kubernetes Orchestrates your application!!

Kubernetes - Background

- Kubernetes came out of [Google](#)! It was open-sourced on 2014 and Handed over to “[Cloud Native Computing Foundation](#) (CNCF).
- Google has been running many of it's systems on containers for years even before Docker came along.
- Google goes through billions of containers per week for applications like Gmail, Google Search , GFS (Google File System) ..etc. Lots of Containers hun??
- Kubernetes is written in Go (Golang)
- Logo – Meaning “the person who steers the ship”
- Kubernetes short name K8S.



Container World Challenges

Application deployment & operational challenges:

- Fail over if one or more nodes experience an outage
- Scale-up & Scale-down: ability to add or remove containers based on application demand
- Zero downtime releases (ZDR)
- Application updates & rollbacks
- Health checks & Self healing systems
 - Networking after scaling and self healing, logs, alerting.. Etc
- Traffic routing and Load Balancing (LB), & More...

What is an Orchestrator?

What is an Orchestrator?

Orchestrator is a system that performs below **without** you having to **supervise**.

- Deploy the application
- Scale it up and down
- Performs rolling updates
- Rollbacks

Kubernetes Advantage:

- All these we can achieve with **declarative** configurations.
- Declare **desired state** of the system rather than executing a **series of instruction**.

Example:

- Declare how many containers, replicas are required for your application (**replicas = 3**)
- It's not about just creating replicas, but it will **continuously monitor** to ensure the desired state.

Orchestrator - Simple Analogy



Foot ball team is made of individuals

- No two are **same**
- Each has **different role** to play in the team
formation and stick to the plan
- Some **defend**, some **attack**, some are great at **passing**
- Some are great at **shooting**
- LW/RW-wingers, LB/RB-backs, AM-attacking mid-fielder

- **Coach** gives everyone a position and organizes them into a **team** with purpose
- Coach makes sure that team maintain the

KUBERNETES – By P Nageswara Rao – Gamut Gurus

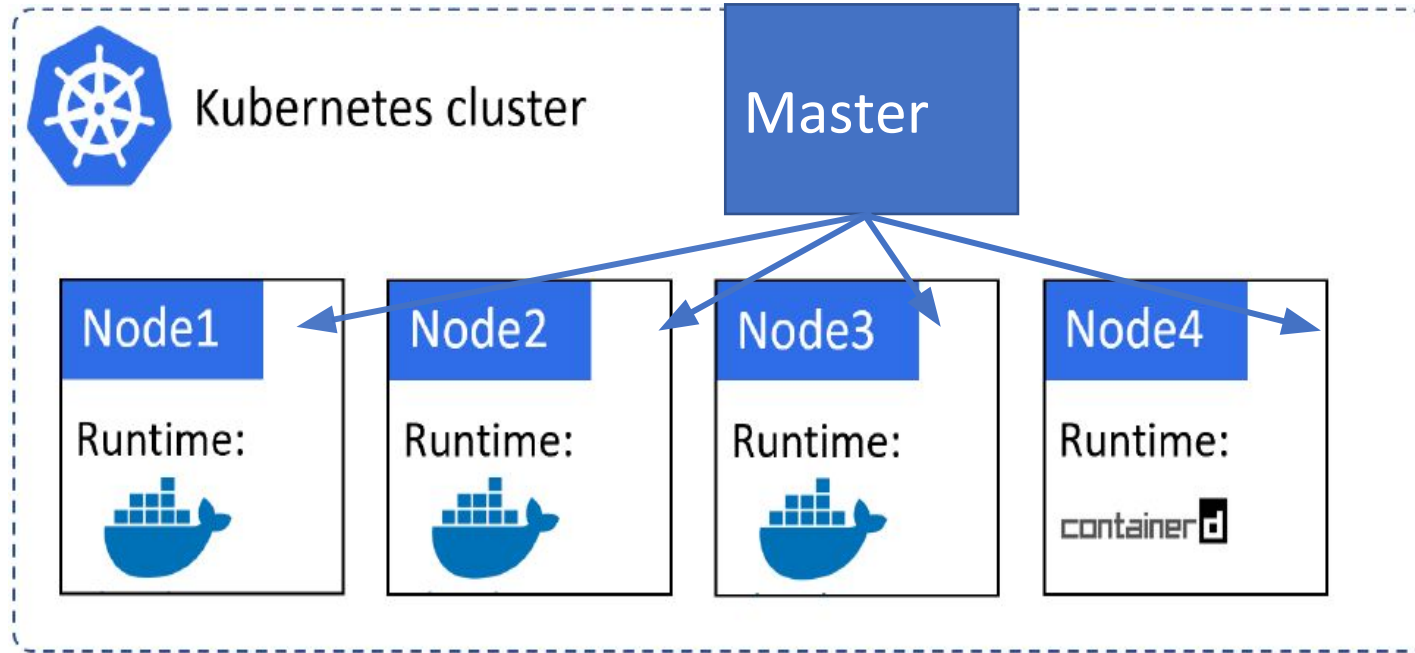
- In the sports world we call this coaching.
 - In the application world we call it orchestration.

In Application World:

- Some serve Web pages
- Some do authentication
- Some do Searches
- Some do store data
- Kubernetes comes along a bit like the coach in the football analogy – and organizes everything into a useful app and keeps things running smoothly.

Kubernetes Cluster

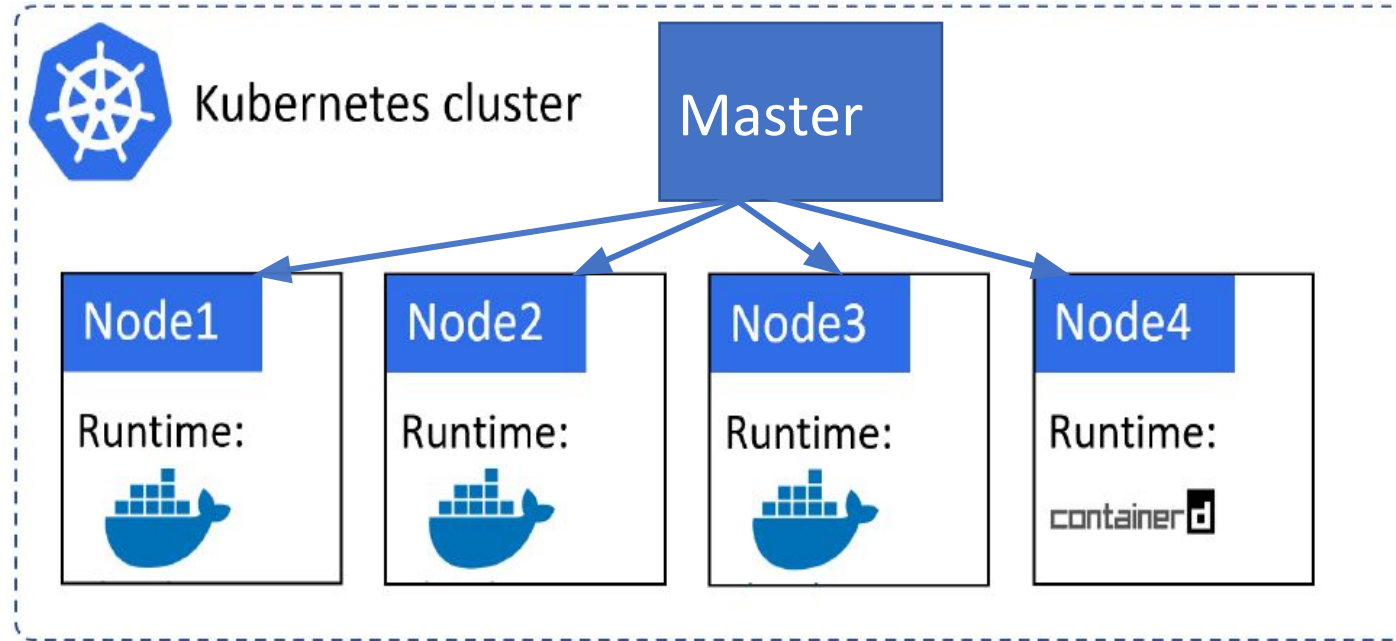
- The **cluster** is made up of one or more **Masters**, and a bunch of **Nodes**.



- You might have K8S cluster with **10 nodes** to run your **production** application.
- Behind the scenes each node is running Docker as it's **container runtime**.
- This means, that Docker is the **low-level technology** that **starts, stops** containers etc.
- **And..** K8S looks after big picture things like.. Deciding **when to scale up or down, update, rollback** ..etc.

Kubernetes Cluster

- Kubernetes **abstracts** the runtime
- Container runtime can be 'Docker' or 'Containerd'
- Containerd is **overtaking** Docker
- Containerd is **stripped-down** version of Docker.
- **Abstraction:-** Facilitated by **Container Runtime Interface** using which you can integrate any **3rd party** container runtime. Thanks CRI!



Kubernetes Vs Docker Swarm

- **Orchestrators:**
 - Docker swarm
 - Mesosphere's DCOS
 - Kubernetes
- During 2016 and 2017 Orchestrator wars, Kubernetes WON with active development and market-share.

Kubernetes - Data centre OS

What is an Data centre OS?

- Sometimes, we call Kubernetes as “Data Center OS”
- In modern Data Centre architectures, we are abandoning traditional view of Data Centre as collection of computers. **Instead..**
- We are viewing it as a single large computer. **What does that mean?**

To understand that in better way.. Look at this example.

- A typical computer is a collection of CPU, RAM, Storage and Networking.
- Having OS, for example.. It's rare for a developer to take care which CPU core or exact Memory address their application uses.
- We let the OS decide all of that! Means it abstracts away all of the CPU, Memory details.

Kubernetes - Data centre OS

- Now, apply this same **abstraction** to data center resources.
- A **typical** Data Center is collection of computers. In **modern** data center architecture, we view the data center as just a **pool of compute, network and storage**.
- This means we **no** longer need to **care about** which server or LUN our containers are running on - just leave this up to the data center OS.
- Gone are the days of taking your app and saying “**Run** this part of the app on this **node**, with this **IP**, on this specific **LUN**...”.
- In the cloud-native Kubernetes world, we’re more about saying “Hey Kubernetes, I’ve got this app and it consists of these parts... **just run it for me please**”.
- Kubernetes then goes off and does all the **hard scheduling** and **orchestration** work.

Data Centre OS - Simple Analogy

Kubernetes is Data Center OS

Think about the process of sending goods via courier services:

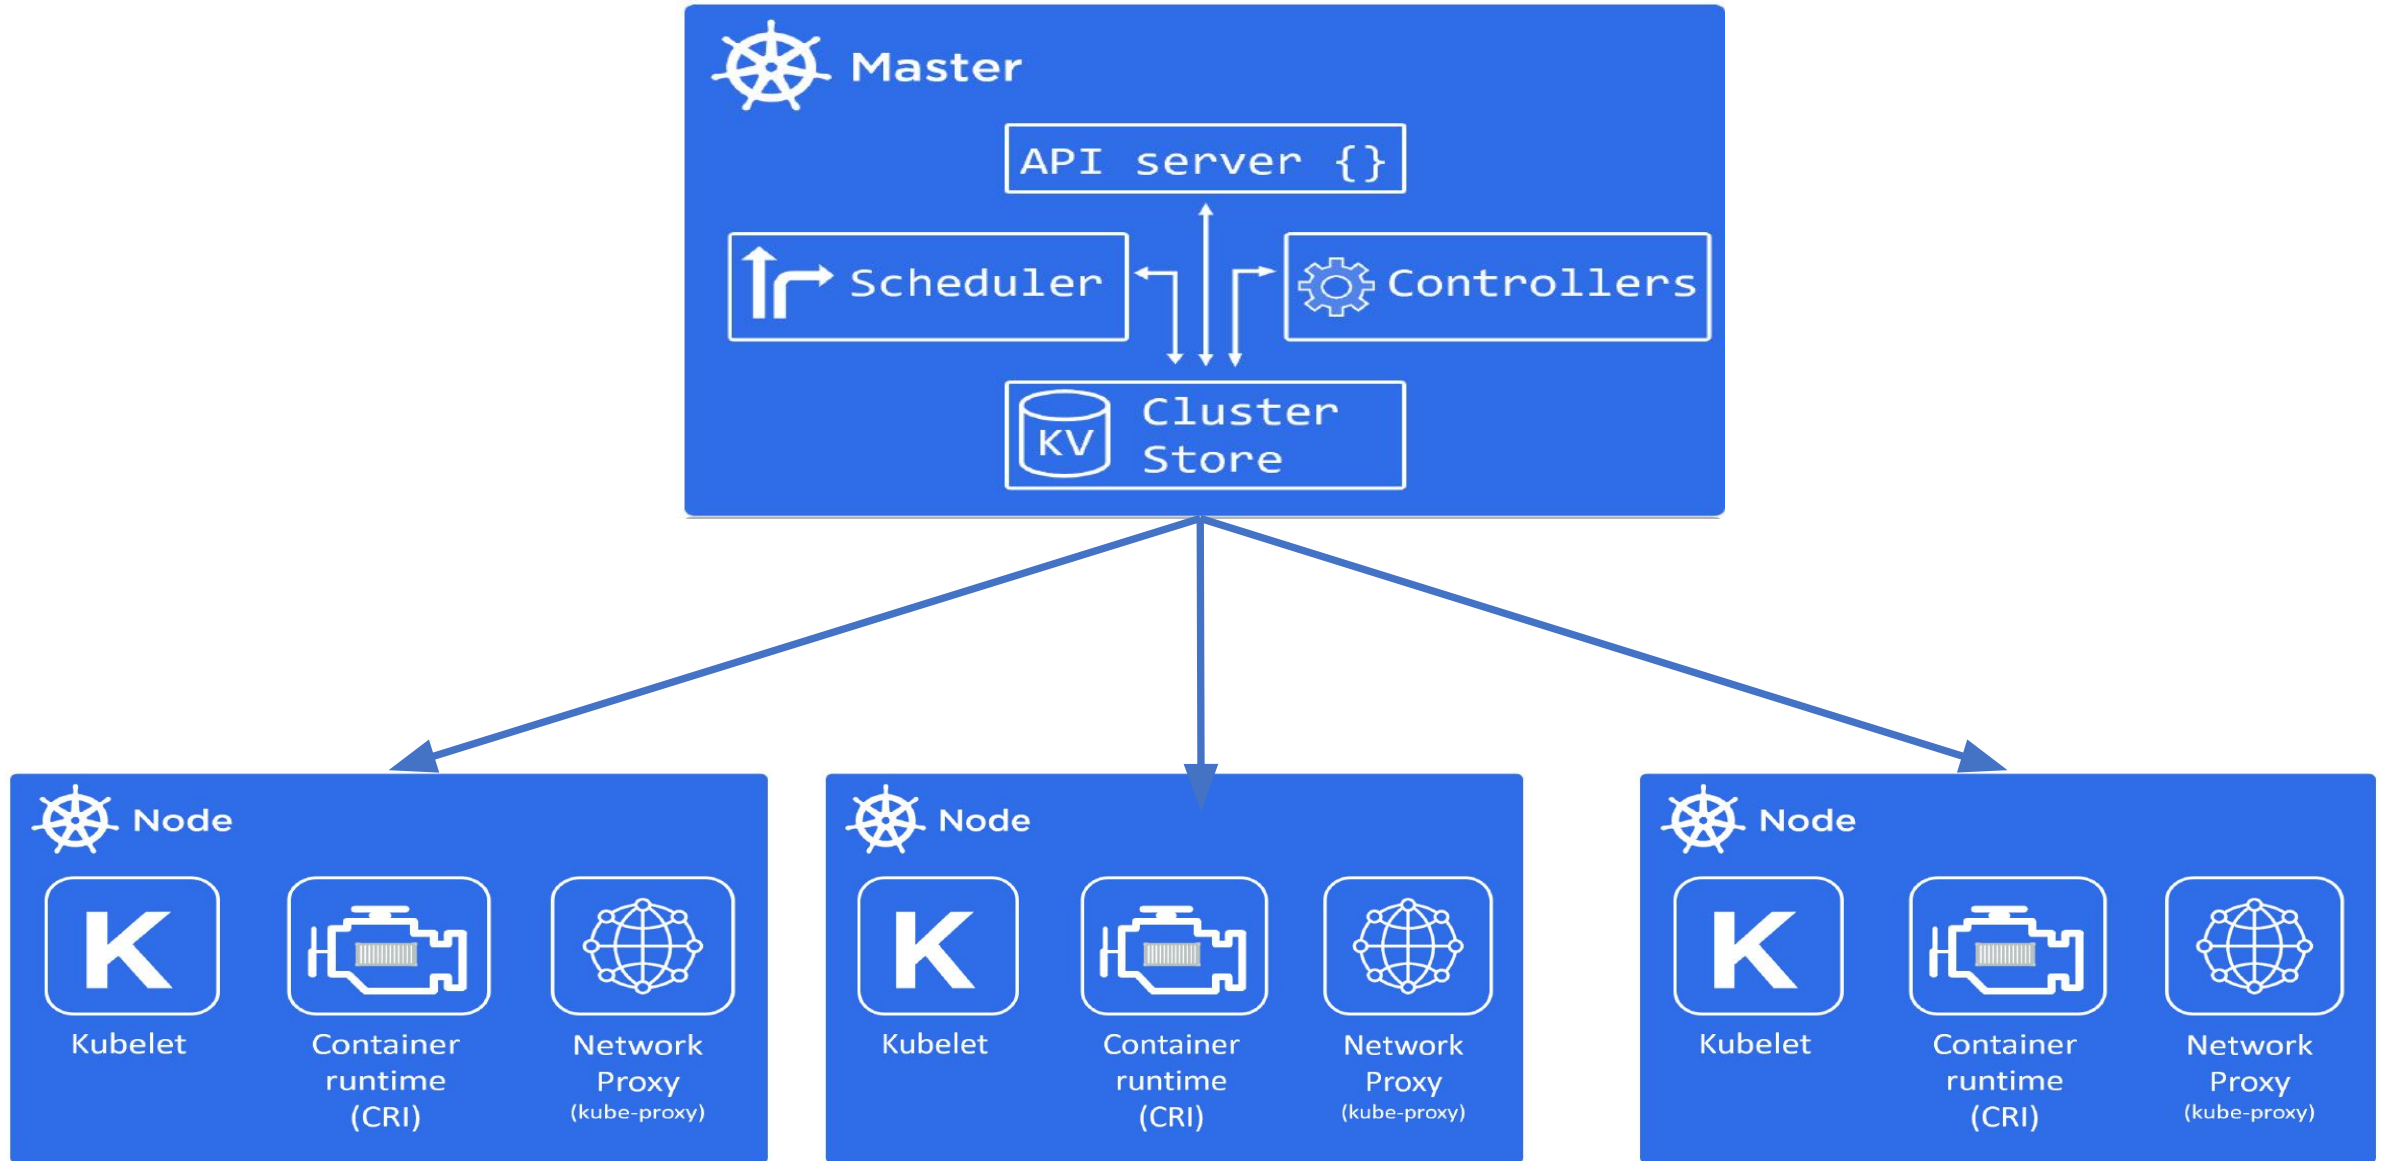
- You need to only **package** and put a **label** on it.
- Rest of the complex logistics like **which** plane, truck, drivers will be taken care by the **courier service** provider
- Kubernetes just **does** the **same** in Data Center



Conclusion

- You **develop** your application with **Docker** and then use **Kubernetes** to **Run**/Orchestrate it.
- Use Kubernetes to **run** your containerized Application in **Test** or **Production** with Orchestrator **capabilities** such as Deploy, Scale-up, Scale-down, Perform updates & rollback without much **supervision**.
- Kubernetes and Docker are **Complimentary** technologies.
- We integrate Docker and Kubernetes to **Package** the application code with dependencies and **Run** the application on Production or Test.
- Kubernetes is a Data centre **OS** (OS for Dev & Kubernetes for Data centre)

Kubernetes Architecture




Kubernetes Architecture

- A Kubernetes cluster is made of **Masters** and **Nodes**
- **Master/Nodes** are **Linux** hosts that can be:
 - **VMs**
 - **Bare** metal servers
 - **Cloud** instances.

Kubernetes Architecture - Master

Master/Control Plane:


- The Master is **in-charge** of the cluster.
- Master:
 - Makes **Scheduling** decisions
 - **Monitors** the cluster
 - **Implement** the Changes
 - Respond to **Events**
- So, we call Master as **Control Plane**.



Kubernetes Architecture - Master

- Control Plane Components:

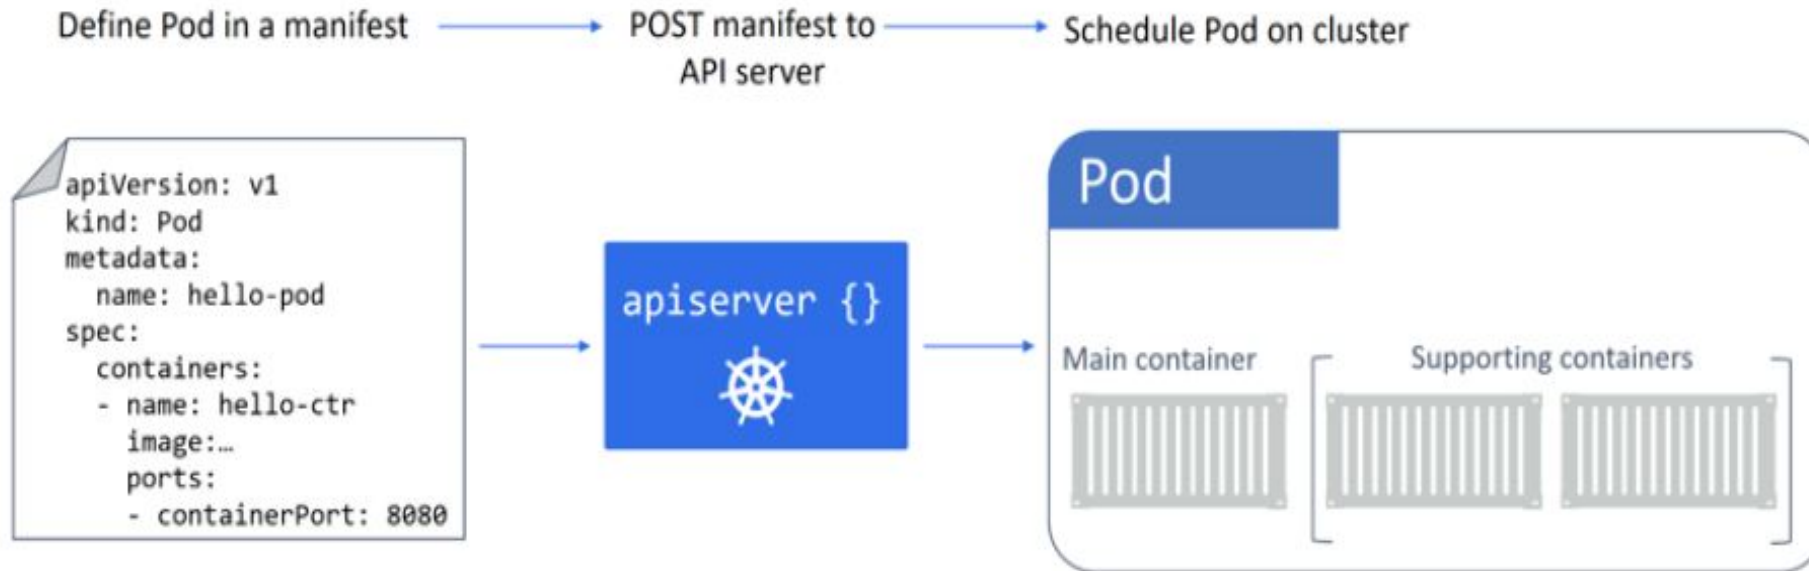
- API Server
- Cluster Store
- Control Manager
- Scheduler



API Server:

- Think of API server as the **brains of the cluster**.
- The API server is **front door** into Kubernetes.
- This is the only component in Control plane we **interact** with.

Kubernetes Architecture - Master




- We **post YAML** configuration files (manifests) to API server.
- This YAML file is **validated**, **persisted** to the **Cluster store** and **deployed** to the cluster.
- These YAML configuration files **contains desired state of the application**.

Example: Which container **image**, Which **ports** to expose, How many **Pod** replicas .. Etc.

Kubernetes Architecture - Master

Cluster Store:


- The API server is **brains** of the cluster and 'Cluster store' is its **memory**.
- It **stores the entire configuration and state** of the cluster.
- **No** cluster store, **No** cluster!
- This cluster store is based on **etcd**, a popular distributed database.



Kubernetes Architecture - Master

Controller Manager:

- It implements several **control loops** that watch the cluster **and** respond to events.
- These Control loops run in **background** mode
- They constantly **watch API** server for changes

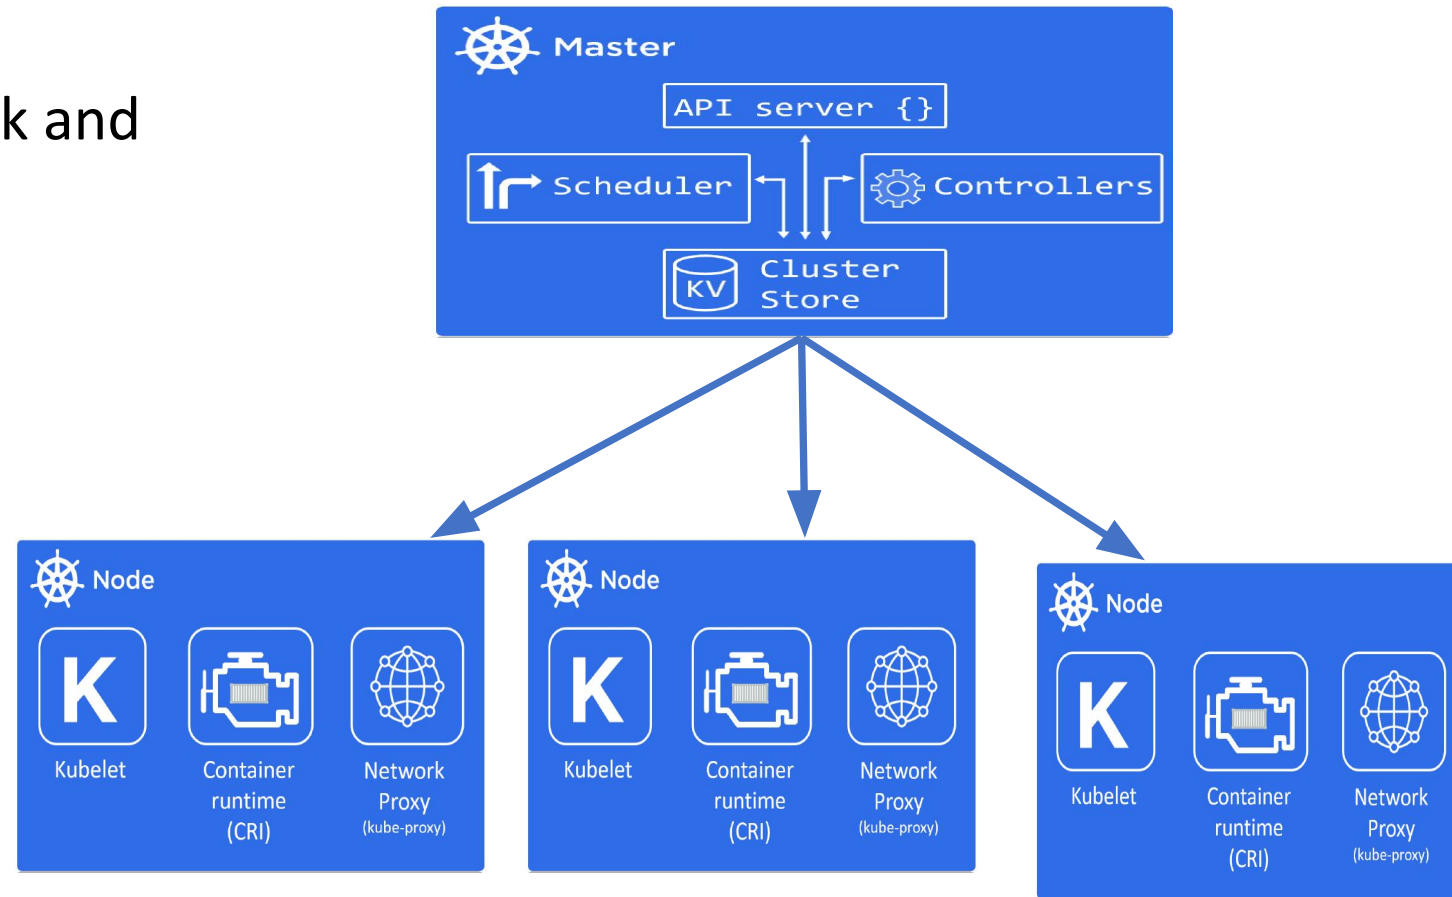


Controller Manager's Aim is to: Ensure the current state of the cluster matches the **desired** state

Kubernetes Architecture - Master

Scheduler:

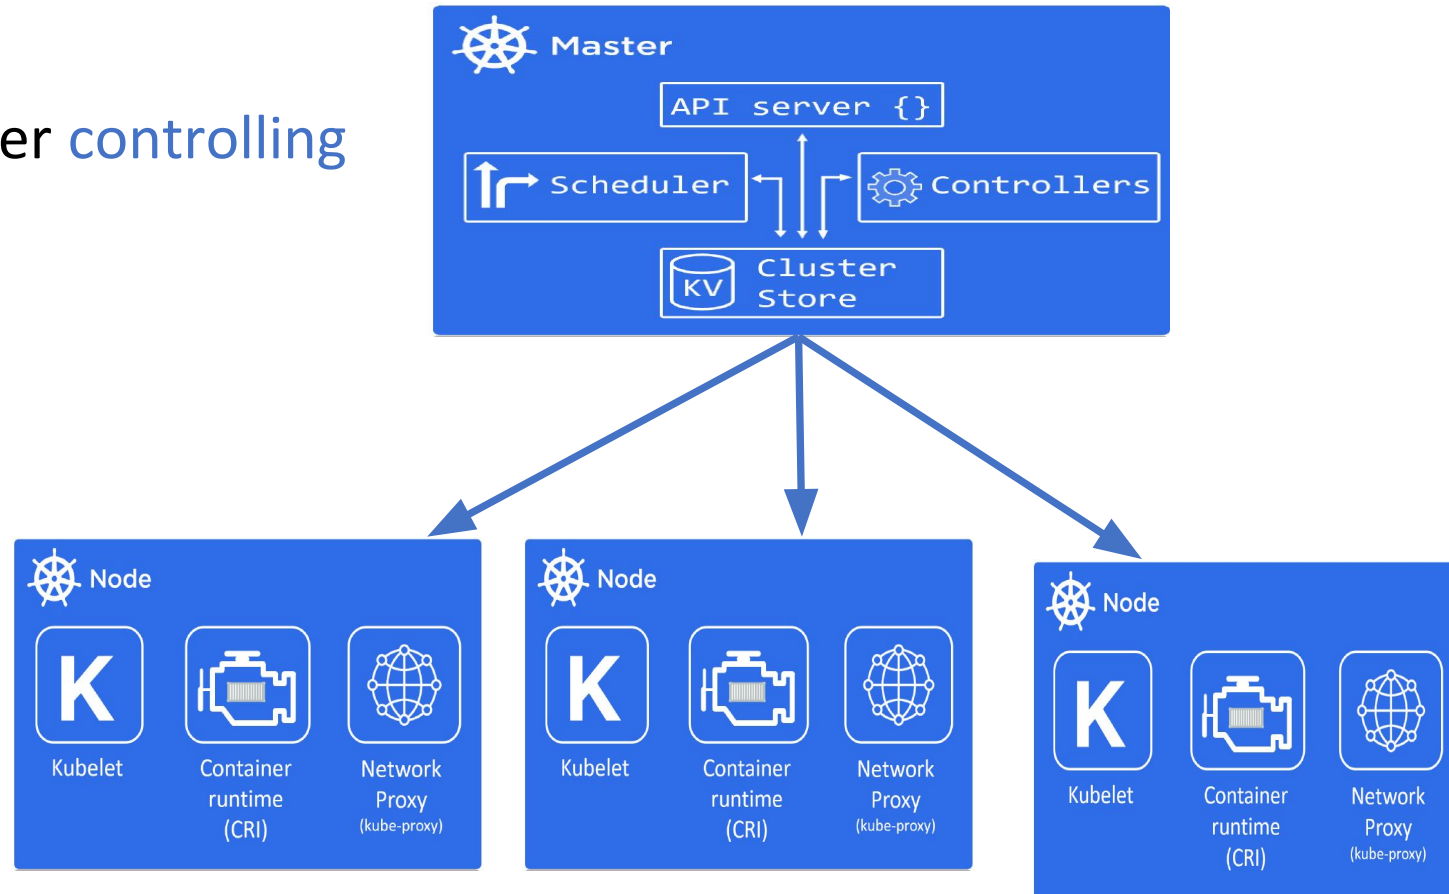
- The scheduler **watches** for new work and assigns it to nodes



Master/Control Plane - Summary

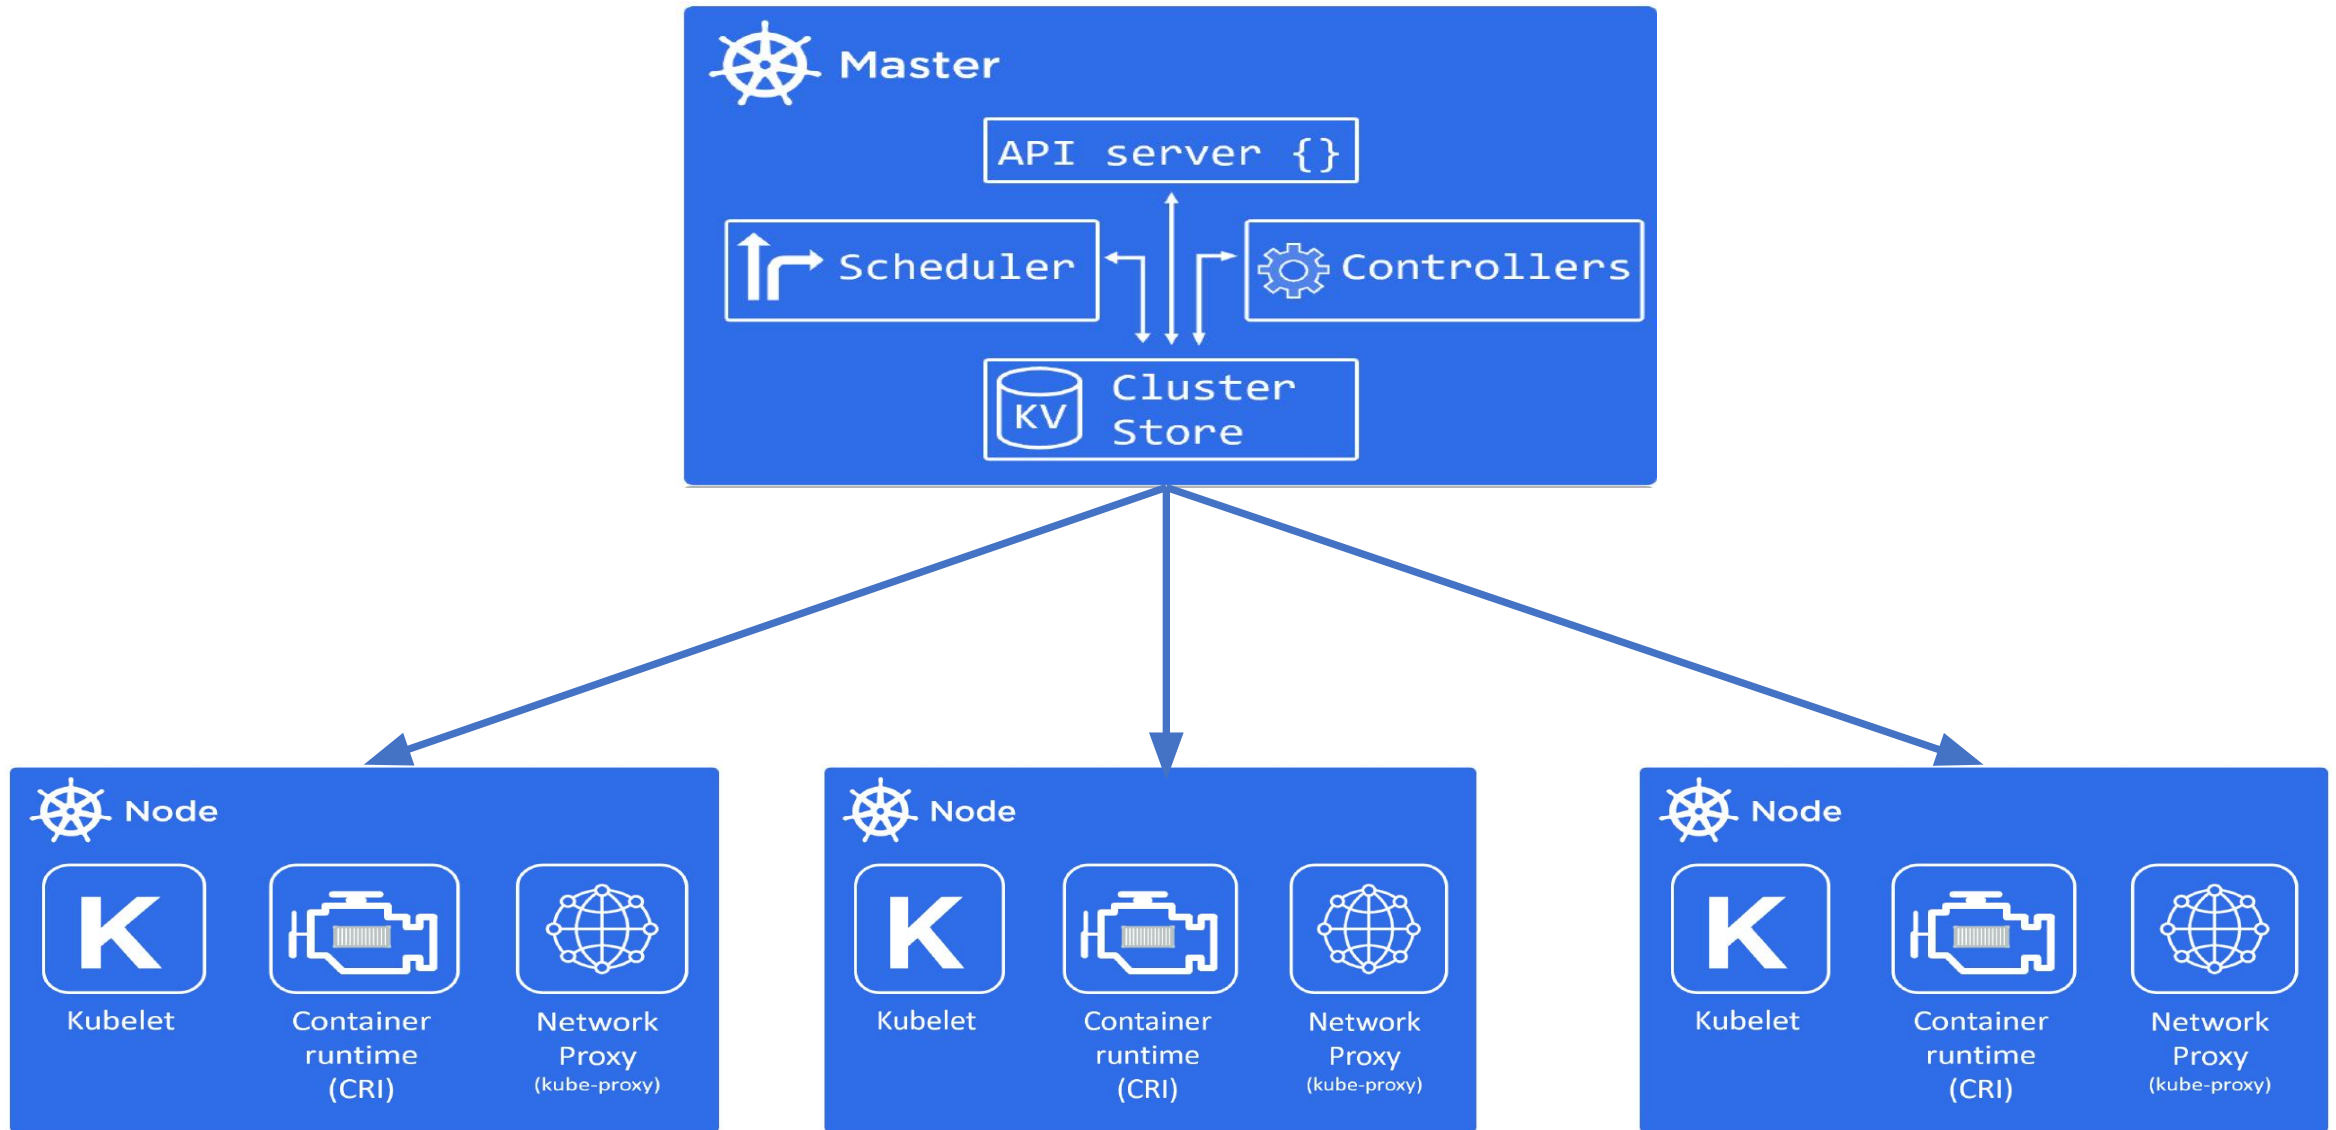
Master Server Summary:

- Master takes care of complete cluster **controlling** and **scheduling**.
- Master components are:
 - API server
 - Scheduler
 - Controllers
 - Cluster store



- It's also considered a good practice **not to run application** workloads on masters.
- This allows masters to **concentrate entirely** on managing the cluster.

Kubernetes Architecture - Nodes



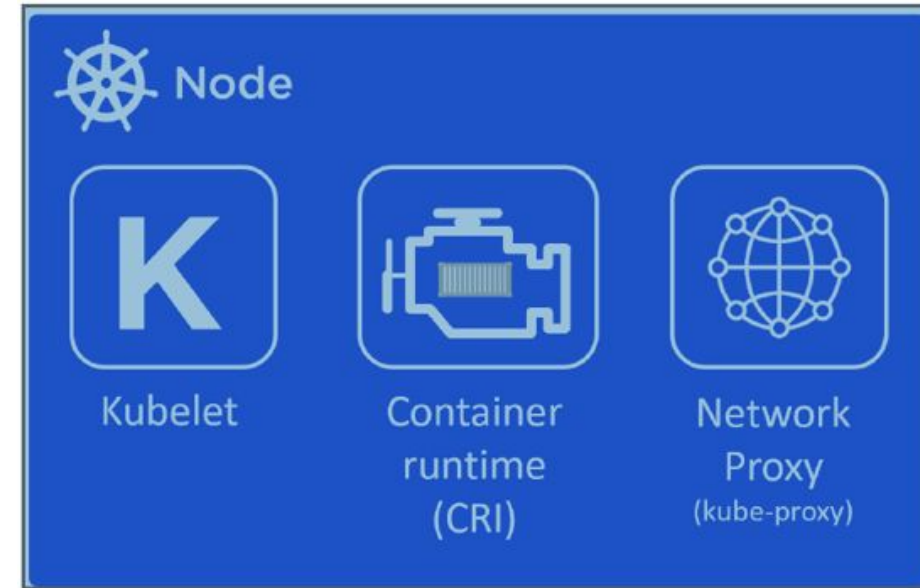
Kubernetes Architecture - Nodes

Nodes/Data Plane:

1. Nodes are the **workers** of a Kubernetes cluster
2. Nodes are where application **services** run.

Nodes do:

- Constantly **watch** the API Server for new work assignments
 - Execute **new work** assignments
 - **Report back** to the control plane
-
- We call Node sometimes as **Data Plane**.



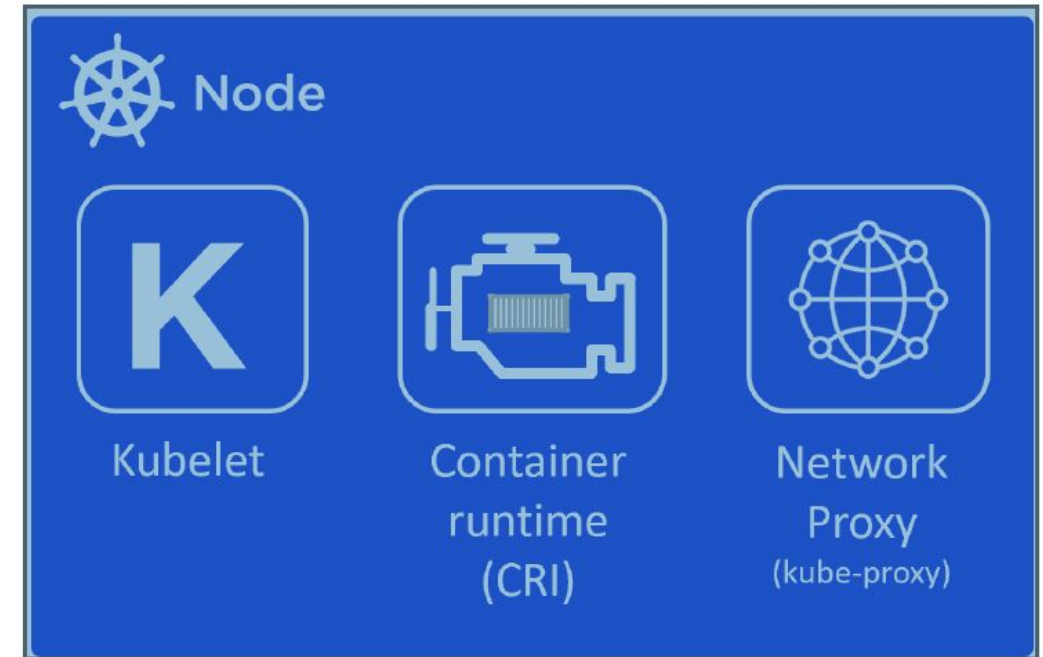
Kubernetes Architecture - Nodes

Node Major Components:

- Kubelet
- Container Runtime
- Kube Proxy

Kubelet:

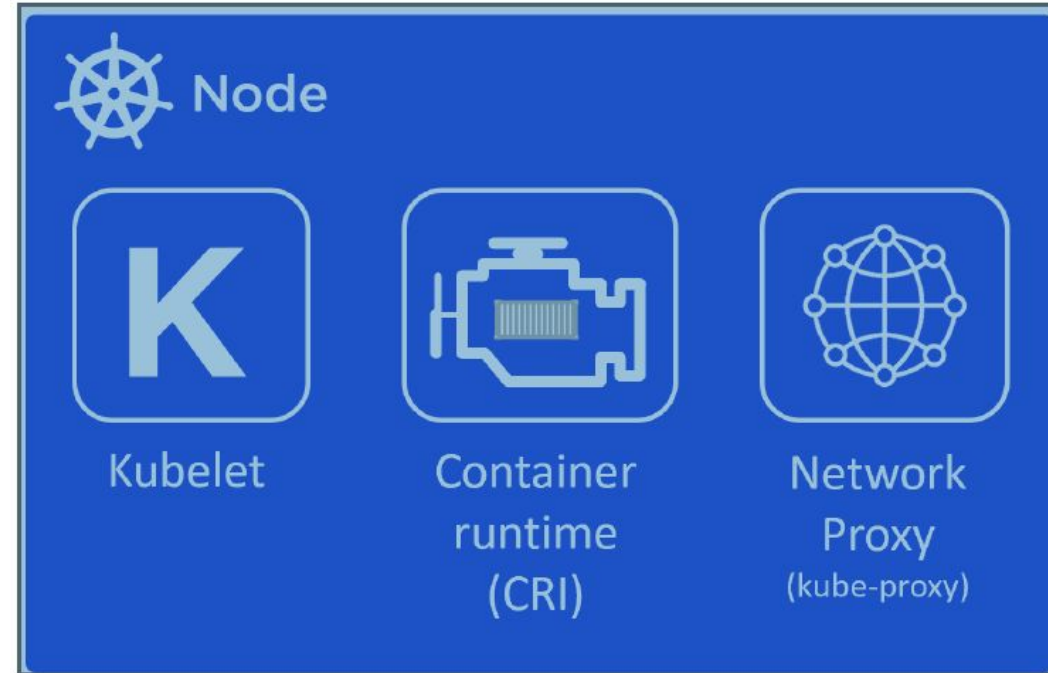
- Kubelet is Kubernetes **agent** that runs on all nodes in the cluster.
- Kubelet is the component that **watches** the API server for **new work** assignments.
- If there is any work, **it carries out the task** and **reports** back to master.
- If **can't** execute any task also, it report back to **master** and master **decides** what next.
- Ex: If a **Pod fails** on a node it's NOT Kubelet's responsibility to find another. But simply reports to master and **master takes care of it**.



Kubernetes Architecture - Nodes

Container Runtime:

- Kubelet need Container runtime for:
 - Pulling the images
 - Starting & Stopping the containers
- Kubernetes supports below container runtimes:
 - Docker
 - Containerd



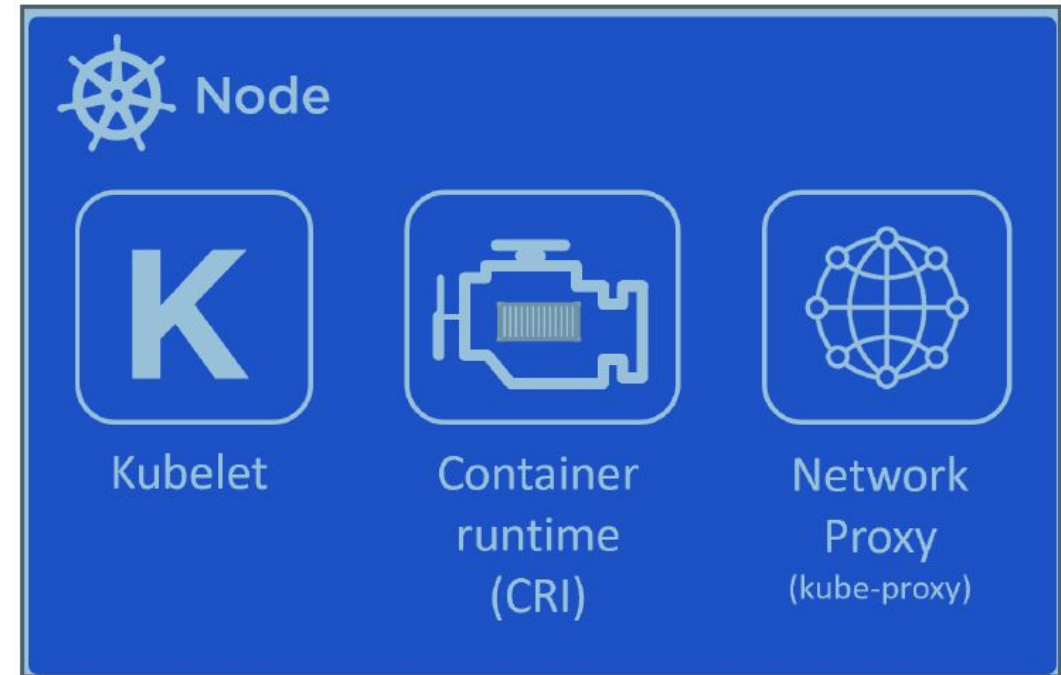
Kubernetes Architecture - Nodes

Kube-proxy:

- Runs on every node of the cluster and is responsible for local networking.

Example:

- It makes sure each node gets it's own unique IP
- Implements IPTables
- IPVS rules to handle Load Balancing



Applications Deployments

Applications Deployments

Microservices


- Docker and Kubernetes are especially useful in **Microservices world**.

There are **two** architectural styles in application software development.

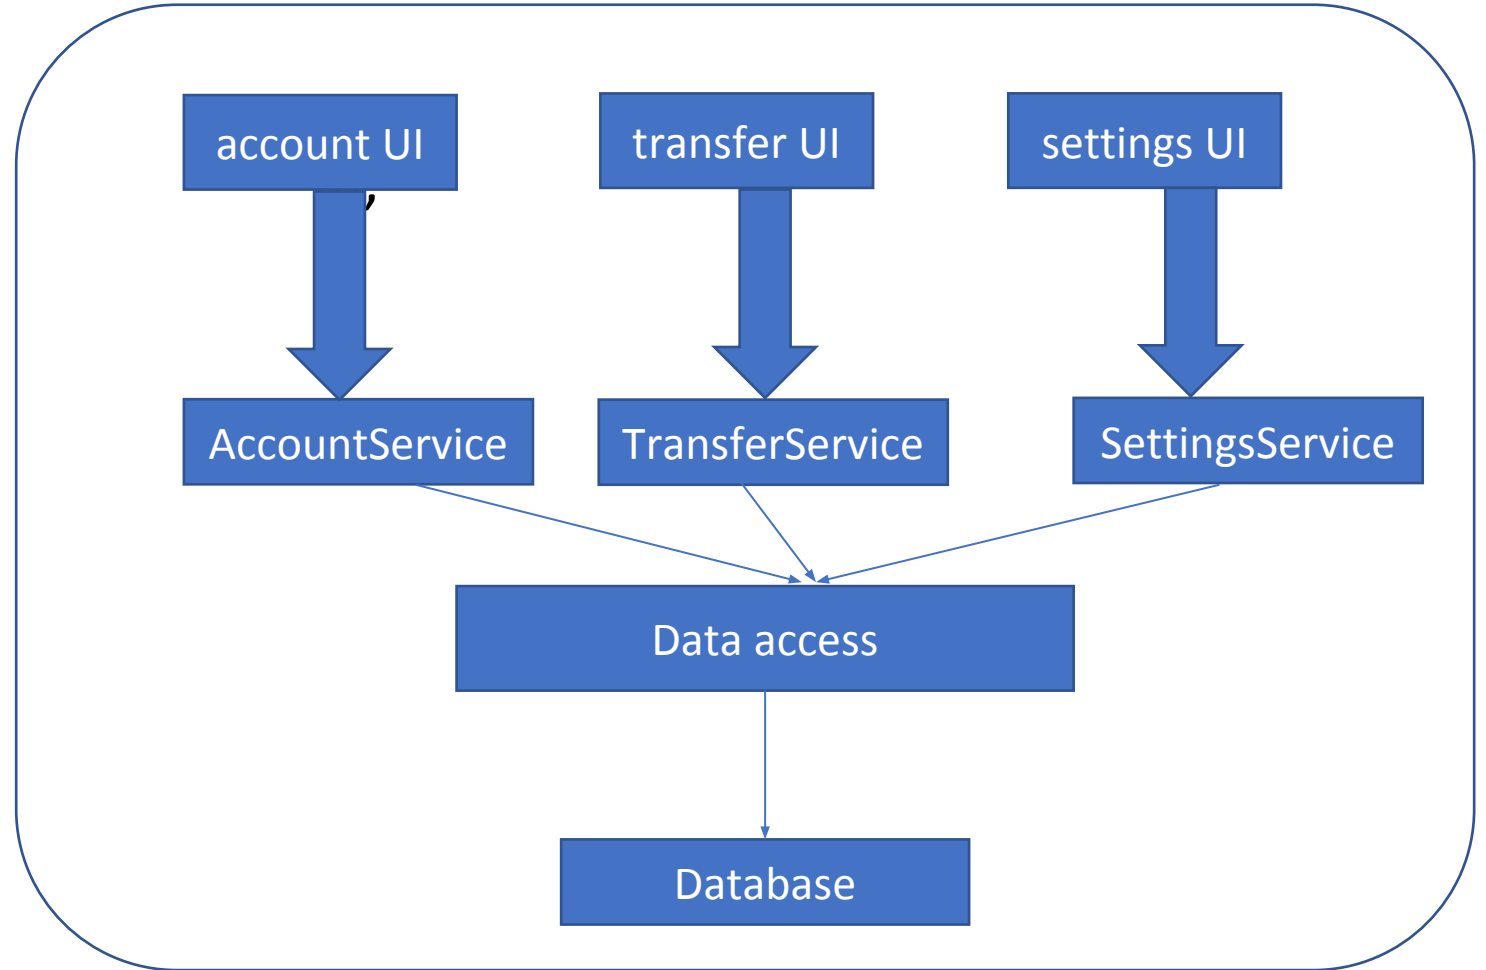
1. Monolithic application
2. Microservices application

- The best way to understand Microservices is to **compare** them with an old, traditional approach for building large applications, the **monolithic** design.

Microservices Vs Monolithic



Monolithic Application



Microservices Application

Microservices

Monolithic Design:

- The whole application is monolith, a **single logical executable**. Good for small projects.
- To make any change to the system, we must build and deploy **whole** application.
- Scaling up is very difficult. What if we want to scale **only** one service ex: DB, Login, Video stream ..etc.
- **Whole** server must be **provisioned** with enough memory and CPU for all. This can be expensive.
- Compile and build times become **longer**. Developers need to **checkout** whole project from SCM, loading in IDE takes more time.
- In case of deployment failure, **whole system is unavailable**.
- Not suitable for **Agile** development and **Continuous integration** processes as we need to release almost at once.

Microservices

Microservices Design:

- This architecture **decomposes** the application into **smaller** pieces.
- **Microservice Architecture(MSA)**, is an architectural style and design pattern which says that an application should consist of a collection of **loosely-coupled** services.
- In other words, **each** of the **services** will have its own **responsibilities** or provide specific functionality, independent of others.
- Each service is deployed on a **separate** host.
- You don't need to **stop whole system** to upgrade a piece of functionality.
- When deployed, microservices improve the **fault tolerance** for entire application.


Ex: if there is **any memory leak** in once service or some other problem, only this service will be **affected** and can then be fixed and upgraded **without interfering** with the rest of the system.

- Each microservice is small, **easier** to understand by developer, easy to develop and deploy.
- Better for **continuous** delivery as small units are easier to manage, test and deploy.

Applications Deployments

Applications Deployments

Running Applications on VM, Container & K8S




Virtualization – Virtual Machine

Docker – Container

Kubernetes – Pod (Pods are just a vehicle for deploying applications)

Pods



- **Virtualization - World**

In the Virtualization world, the **atomic unit** of scheduling is the **Virtual Machine**.

That means, We deploy the applications on VMs

- **Docker – World**

In The Docker world, the atomic unit is the **Container**.

This means , We deploy the applications inside of containers.


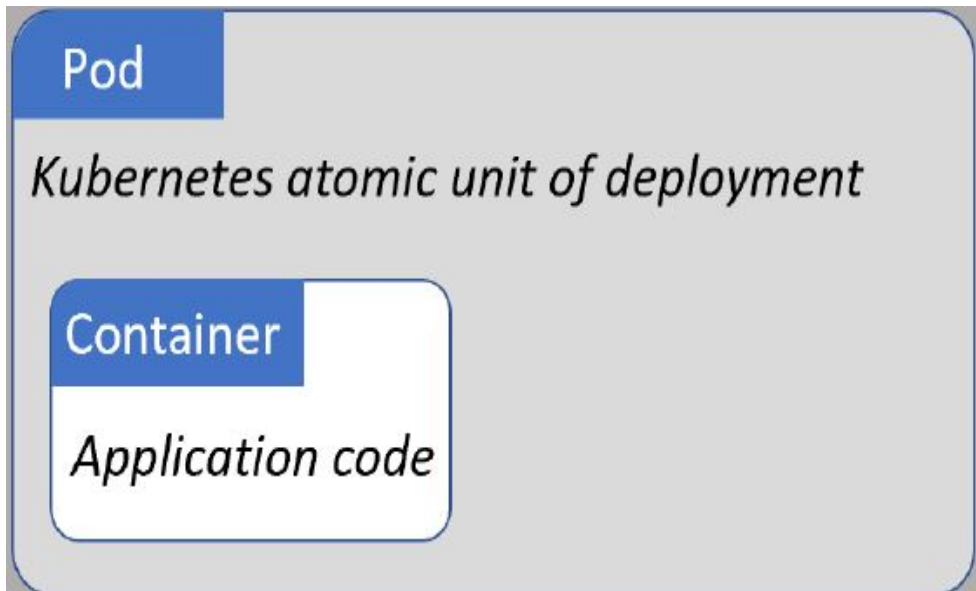
- **Kubernetes – World**

In the Kubernetes world, the atomic unit is the **Pod**.

This means, We deploy the applications in Pods

POD

- We **can't run** a container **directly** on a Kubernetes cluster
- Containers must always **run inside of Pods!**
- Pod can have **single** or **multiple** containers. All container **share** same Pod resources.
- Pods connect to each other using Pod's '**localhost**' network interface.



Namespace & Cgroups

Docker is built based on below Linux Kernel features:

- **NameSpace**
 - **Cgroups**
- 'Namespace' and 'Cgroups' isolates the processes from each other.

NameSpace: Isolates processes, networking, file system ..etc.

Cgroups: Limits how much CPU, Memory one of the process/container is using.


Namespace & Cgroups

What is POD:

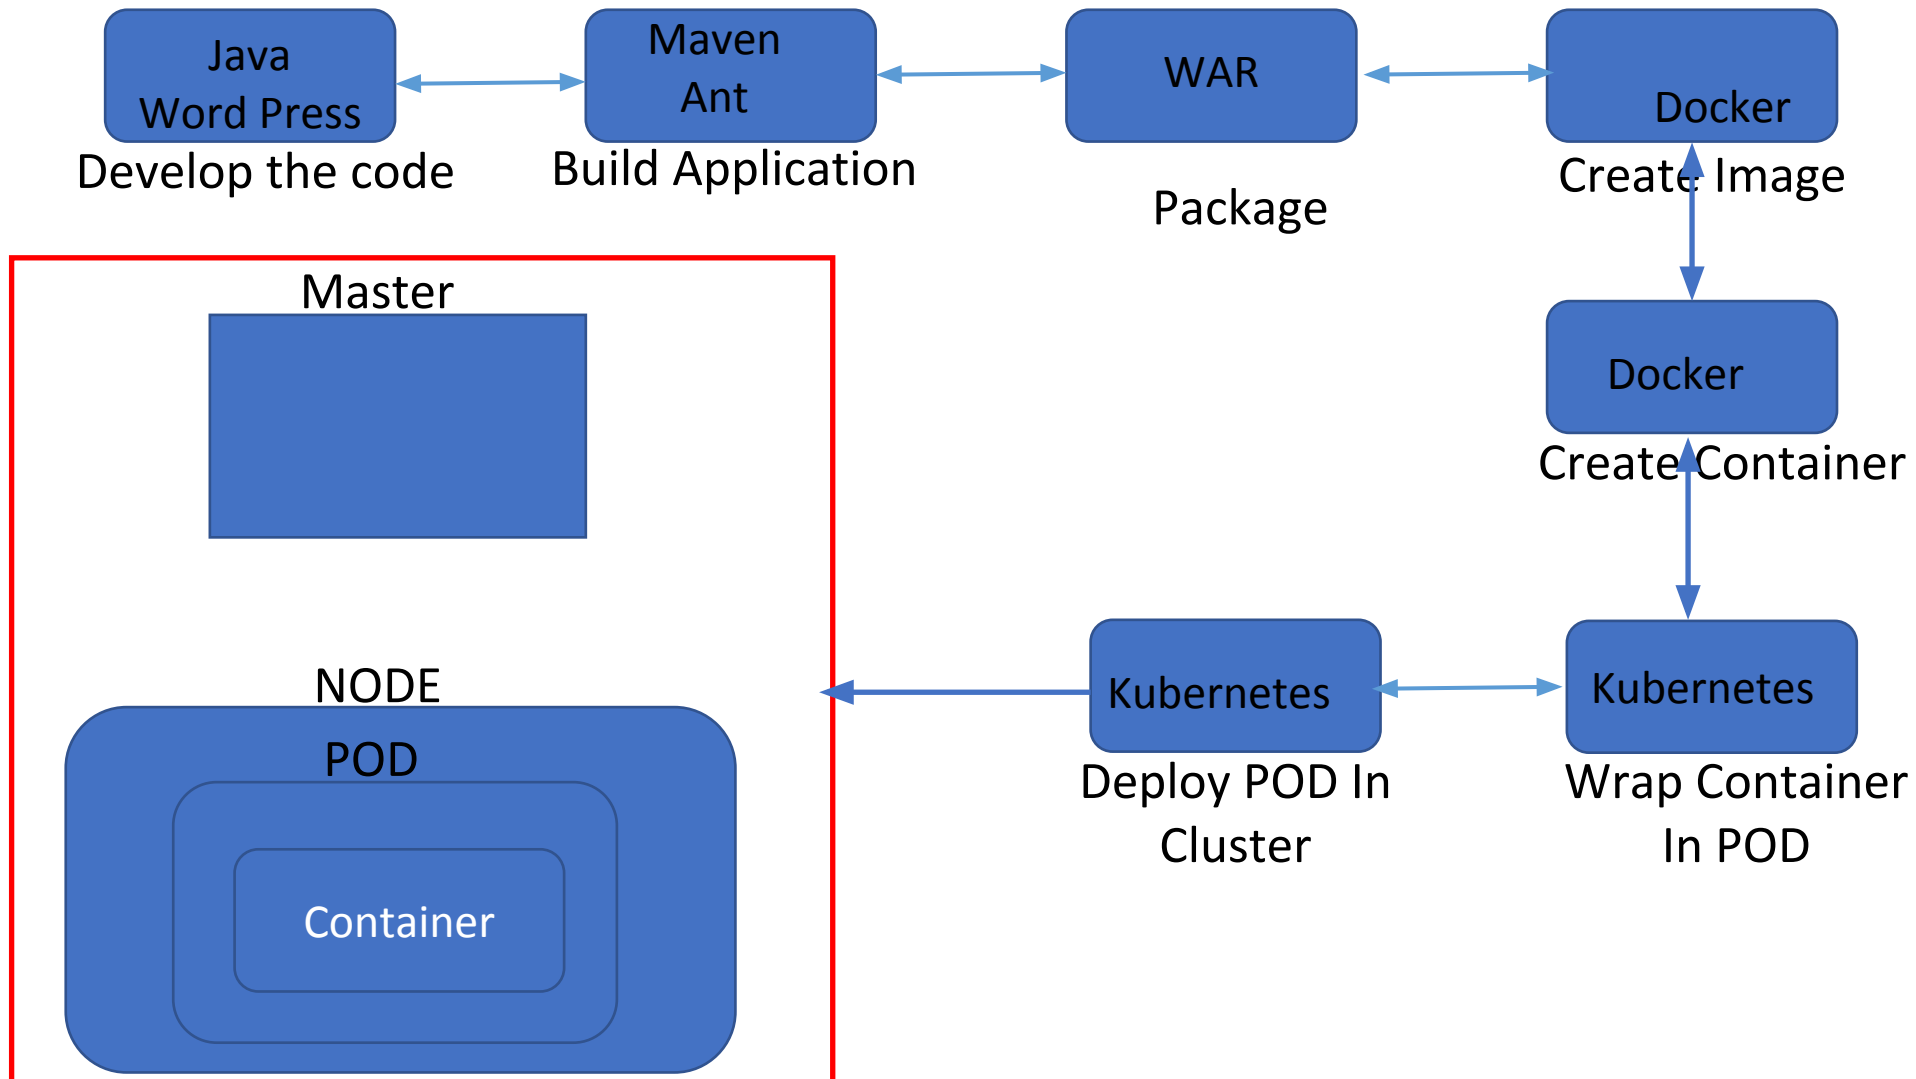
- A Pod is a **shared** execution environment for one or more containers.
- These include: **IP** address, **Ports**, **Hostname**, **Memory**, **Volumes** ..etc.

Pods – Atomic Unit

- The deployment of Pod is an **atomic** operation.
- If we have to **scale** application, **we add or remove Pods**.
- We **don't** scale by **adding** more **containers** to an existing Pod.



Applications Deployments In Kubernetes



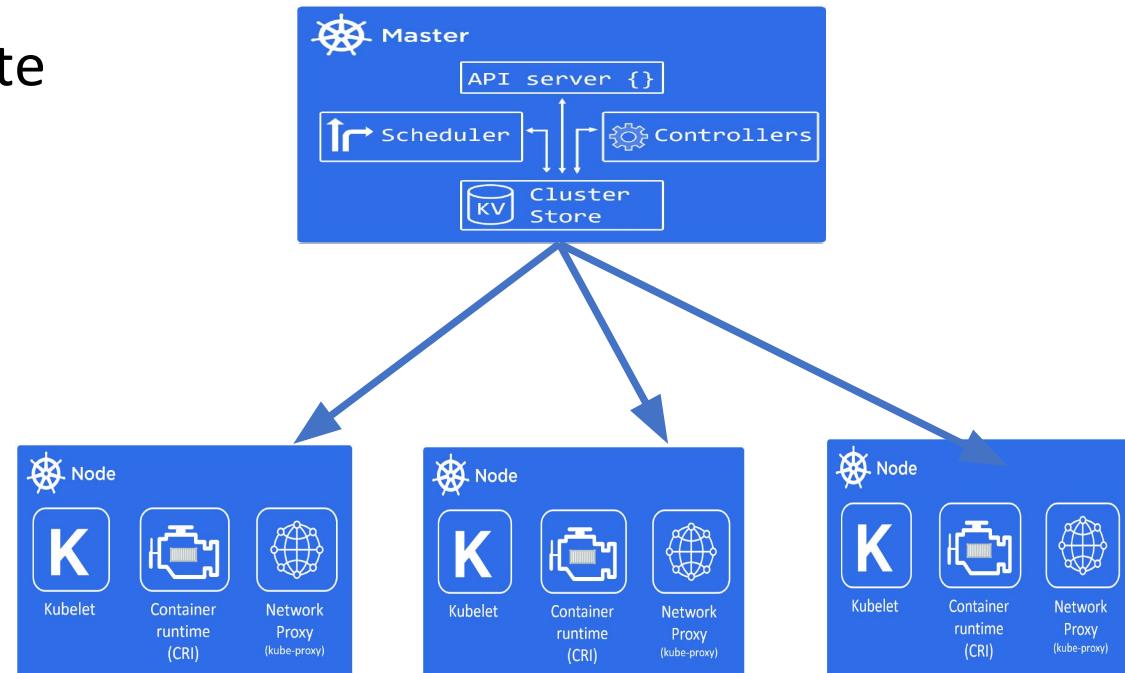
Running Applications on Kubernetes Cluster

Summary:

1. Write the application as small independent services
2. Package each service in it's own container
3. Wrap each container in it's own Pod
4. Deploy the Pods to the cluster

Running Applications on Kubernetes Cluster

1. We **declare** the **desired state** of an application in a manifest file
2. We **POST** it to the Kubernetes **API** server
3. Kubernetes **stores** this in the **cluster store** as the application's desired state
4. Kubernetes **implements** the **desired state** on the cluster
5. Kubernetes **implements watch loops** to make sure the current state of the application doesn't vary from desired state



What's Next?

Pods don't:

self-heal

scale

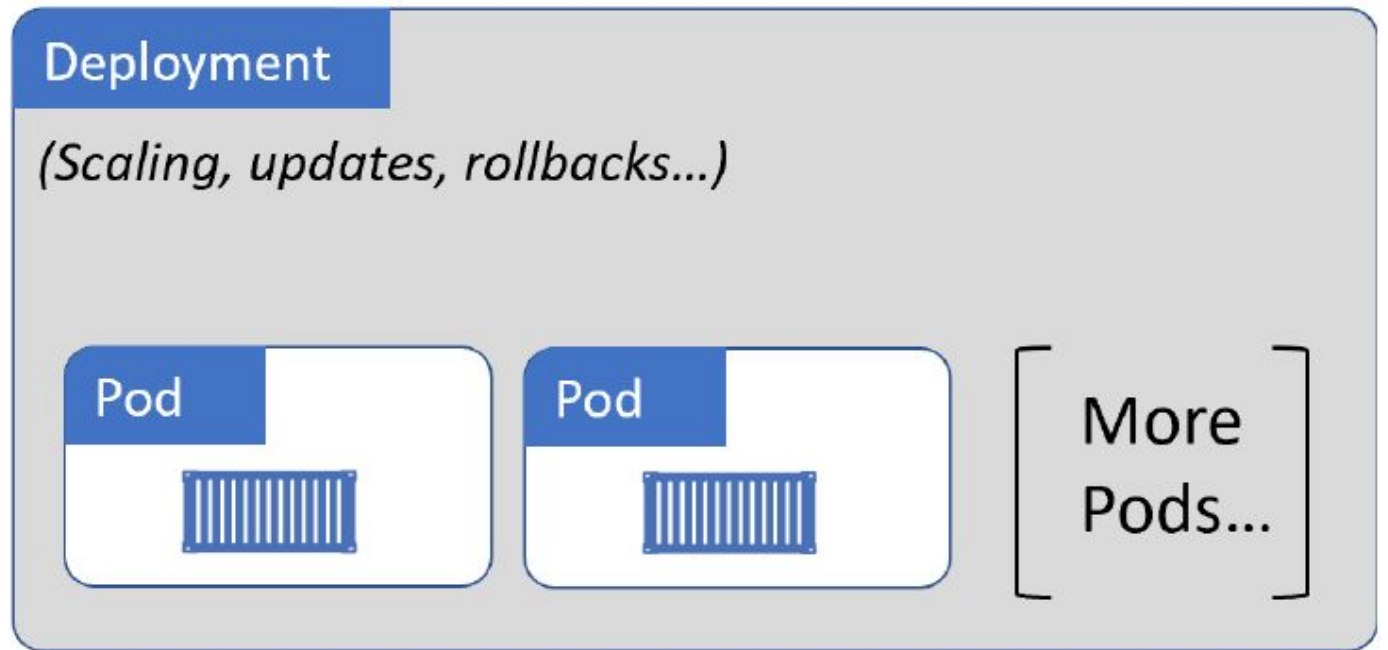
Allow for easy updates

Deployments

- Pods don't self-heal, they don't scale, and they don't allow for easy updates.

\$ ENTER DEPLOYMENTS..

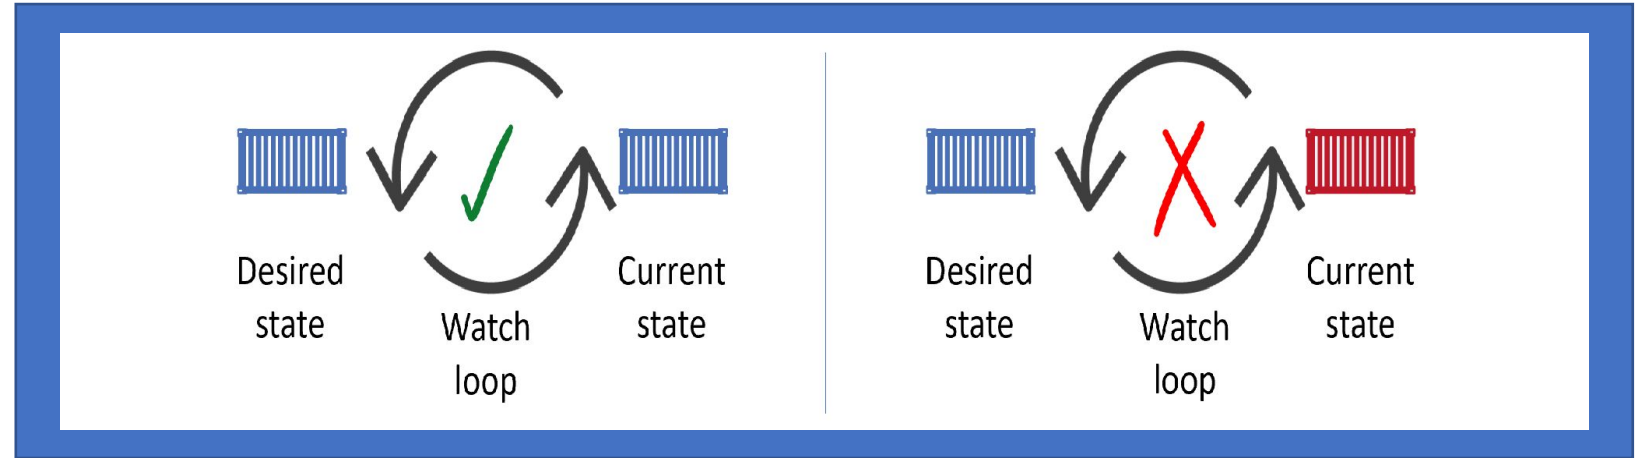
- Deployments do all things like
 - "scale" (scale-up/down)
 - "self-heal"
 - "rolling updates"
 - "roll backs"



- That's why we almost always deploy Pods via 'Deployments'

Desired State.. Current State.. Declarative Model

Just a Concept:



- *Desired State: is What we want*
- *Current State: is What you have*
- *Declarative Model: is a way of telling Kubernetes what our desired state is, without getting into the detail of how to implement it.*

Desired State.. Current Stage..Declarative Model

“Self-healing” with “Deployments”– Deep Dive:

If Pod being managed by Deployment fails, It will be replaced.

Example:

- You need 5 FE and 2 BE Pods.
- Just declare desired state. Sit back! K8S does all the hard work.

Not just that..

- Kubernetes also implements **watch loops** to make sure current state matches desired state.
- If 2 FE Pods go down, Kubernetes automatically detects the discrepancy and brings 2 Pods back up.

With this, you can sleep peacefully at 4:30 AM 😊